McDonalds_Menu - Linear Regression

March 3, 2018

1 McDonald's Menu Dataset

The data consists of nutrients value of all the items that are served @McDonald's Outlets. The dataset is taken from kaggle.

2 Load data

```
In [1]: import numpy as np
        import pandas as pd
        data = pd.read_csv("C:/Users/Aravind/Documents/menu.csv")
        data.head()
```

Out[1]:

	Category	Item	Serving Size	Calories	\
0	Breakfast	Egg McMuffin	4.8 oz (136 g)	300	
1	Breakfast	Egg White Delight	4.8 oz (135 g)	250	
2	Breakfast	Sausage McMuffin	3.9 oz (111 g)	370	
3	Breakfast	Sausage McMuffin with Egg	5.7 oz (161 g)	450	
4	Breakfast	Sausage McMuffin with Egg Whites	5.7 oz (161 g)	400	

	Calories from Fat	Total Fat	Total Fat (% Daily Value)	Saturated Fat	\
0	120	13.0	20	5.0	
1	70	8.0	12	3.0	
2	200	23.0	35	8.0	
3	250	28.0	43	10.0	
4	210	23.0	35	8.0	

	Saturated Fat (% Daily Value)	Trans Fat	...	\
0	25	0.0	...	
1	15	0.0	...	
2	42	0.0	...	
3	52	0.0	...	
4	42	0.0	...	

	Carbohydrates	Carbohydrates (% Daily Value)	Dietary Fiber	\
0	31	10	4	
1	30	10	4	

```
2              29                              10               4
3              30                              10               4
4              30                              10               4

   Dietary Fiber (% Daily Value)  Sugars  Protein  Vitamin A (% Daily Value)  \
0                             17       3       17                         10
1                             17       3       18                          6
2                             17       2       14                          8
3                             17       2       21                         15
4                             17       2       21                          6

   Vitamin C (% Daily Value)  Calcium (% Daily Value)  Iron (% Daily Value)
0                          0                       25                    15
1                          0                       25                     8
2                          0                       25                    10
3                          0                       30                    15
4                          0                       25                    10

[5 rows x 24 columns]
```

```
In [2]: data.shape
```

```
Out[2]: (260, 24)
```

There are 260 rows and 24 columns # Check for missing values

```
In [3]: data.isnull().sum()
```

```
Out[3]: Category                      0
        Item                          0
        Serving Size                  0
        Calories                      0
        Calories from Fat             0
        Total Fat                     0
        Total Fat (% Daily Value)     0
        Saturated Fat                 0
        Saturated Fat (% Daily Value) 0
        Trans Fat                     0
        Cholesterol                   0
        Cholesterol (% Daily Value)   0
        Sodium                        0
        Sodium (% Daily Value)        0
        Carbohydrates                 0
        Carbohydrates (% Daily Value) 0
        Dietary Fiber                 0
        Dietary Fiber (% Daily Value) 0
        Sugars                        0
        Protein                       0
        Vitamin A (% Daily Value)     0
```

```
Vitamin C (% Daily Value)        0
Calcium (% Daily Value)          0
Iron (% Daily Value)             0
dtype: int64
```

There are no missing values # Describe data

```
In [4]: data.describe()
```

Out[4]:

	Calories	Calories from Fat	Total Fat	Total Fat (% Daily Value)
count	260.000000	260.000000	260.000000	260.000000
mean	368.269231	127.096154	14.165385	21.815385
std	240.269886	127.875914	14.205998	21.885199
min	0.000000	0.000000	0.000000	0.000000
25%	210.000000	20.000000	2.375000	3.750000
50%	340.000000	100.000000	11.000000	17.000000
75%	500.000000	200.000000	22.250000	35.000000
max	1880.000000	1060.000000	118.000000	182.000000

	Saturated Fat	Saturated Fat (% Daily Value)	Trans Fat	Cholesterol
count	260.000000	260.000000	260.000000	260.000000
mean	6.007692	29.965385	0.203846	54.942308
std	5.321873	26.639209	0.429133	87.269257
min	0.000000	0.000000	0.000000	0.000000
25%	1.000000	4.750000	0.000000	5.000000
50%	5.000000	24.000000	0.000000	35.000000
75%	10.000000	48.000000	0.000000	65.000000
max	20.000000	102.000000	2.500000	575.000000

	Cholesterol (% Daily Value)	Sodium	...
count	260.000000	260.000000	...
mean	18.392308	495.750000	...
std	29.091653	577.026323	...
min	0.000000	0.000000	...
25%	2.000000	107.500000	...
50%	11.000000	190.000000	...
75%	21.250000	865.000000	...
max	192.000000	3600.000000	...

	Carbohydrates	Carbohydrates (% Daily Value)	Dietary Fiber
count	260.000000	260.000000	260.000000
mean	47.346154	15.780769	1.630769
std	28.252232	9.419544	1.567717
min	0.000000	0.000000	0.000000
25%	30.000000	10.000000	0.000000
50%	44.000000	15.000000	1.000000
75%	60.000000	20.000000	3.000000
max	141.000000	47.000000	7.000000

```
       Dietary Fiber (% Daily Value)      Sugars     Protein  \
count                     260.000000  260.000000  260.000000
mean                        6.530769   29.423077   13.338462
std                         6.307057   28.679797   11.426146
min                         0.000000    0.000000    0.000000
25%                         0.000000    5.750000    4.000000
50%                         5.000000   17.500000   12.000000
75%                        10.000000   48.000000   19.000000
max                        28.000000  128.000000   87.000000

       Vitamin A (% Daily Value)  Vitamin C (% Daily Value)  \
count                 260.000000                 260.000000
mean                   13.426923                   8.534615
std                    24.366381                  26.345542
min                     0.000000                   0.000000
25%                     2.000000                   0.000000
50%                     8.000000                   0.000000
75%                    15.000000                   4.000000
max                   170.000000                 240.000000

       Calcium (% Daily Value)  Iron (% Daily Value)
count               260.000000            260.000000
mean                 20.973077              7.734615
std                  17.019953              8.723263
min                   0.000000              0.000000
25%                   6.000000              0.000000
50%                  20.000000              4.000000
75%                  30.000000             15.000000
max                  70.000000             40.000000

[8 rows x 21 columns]
```

Some variables with daily values and item size, category, etc are not necessary for us. lets drop them.

```
In [5]: dataset = data.loc[:,['Calories','Total Fat','Cholesterol','Sodium','Carbohydrates','Di
        dataset.head()
```

```
Out[5]:    Calories  Total Fat  Cholesterol  Sodium  Carbohydrates  Dietary Fiber  \
        0       300       13.0          260     750             31              4
        1       250        8.0           25     770             30              4
        2       370       23.0           45     780             29              4
        3       450       28.0          285     860             30              4
        4       400       23.0           50     880             30              4

           Sugars  Protein
        0       3       17
```

1	3	18
2	2	14
3	2	21
4	2	21

3 Visualization

Scatter plot

```
In [6]: import matplotlib as mpl
        import matplotlib.pyplot as plt
        import seaborn as sns
        %matplotlib inline
        sns.pairplot(dataset)
```

```
Out[6]: <seaborn.axisgrid.PairGrid at 0x1ff60bbf98>
```

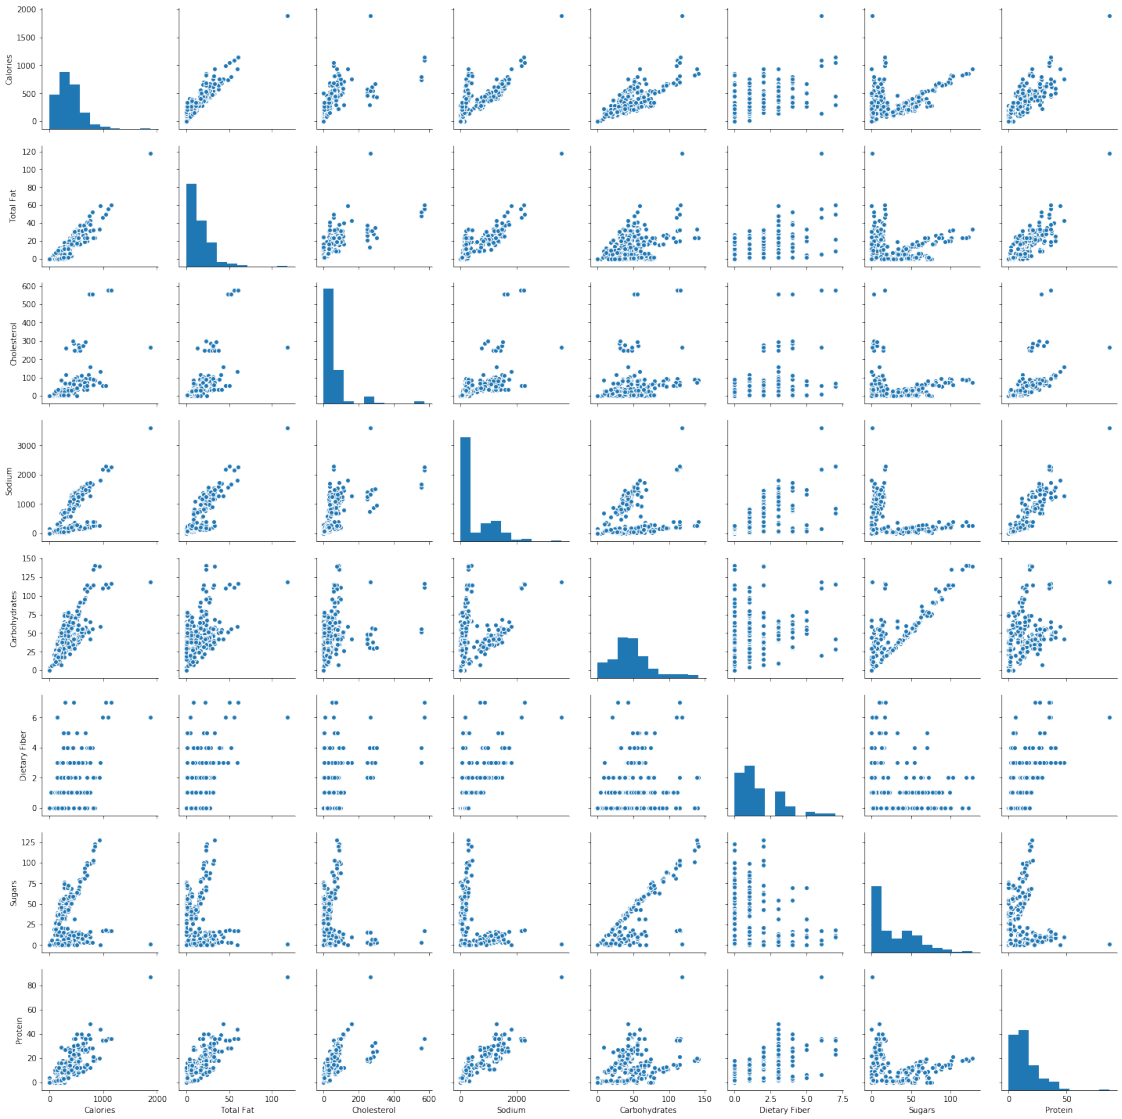

```
In [43]: sns.pairplot(dataset, x_vars=['Total Fat','Carbohydrates','Protein'],y_vars='Calories
```

```
Out[43]: <seaborn.axisgrid.PairGrid at 0x4af29ab780>
```

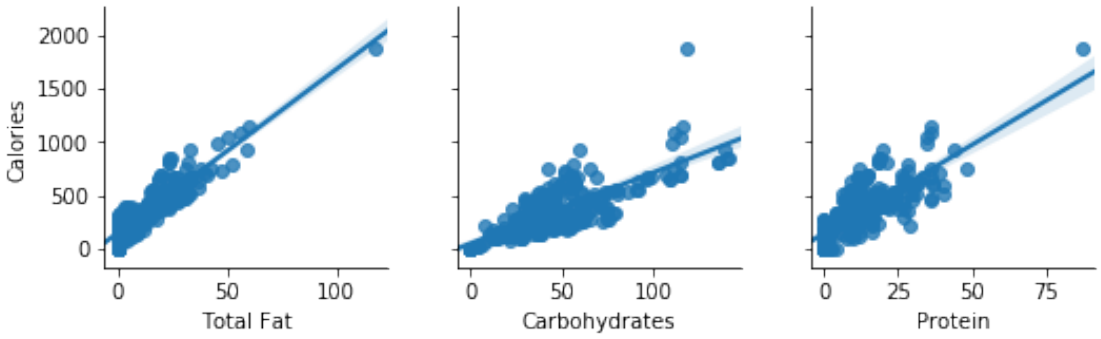

From scatter plots we can see linear relationships between calories and other variable Lets see corrplots

```
In [7]: cor = dataset.corr()
        print(cor)
```

```
                Calories  Total Fat  Cholesterol    Sodium  Carbohydrates  \
Calories        1.000000   0.904409     0.596399  0.712309       0.781539
Total Fat       0.904409   1.000000     0.680547  0.846158       0.461213
Cholesterol     0.596399   0.680547     1.000000  0.624362       0.270977
Sodium          0.712309   0.846158     0.624362  1.000000       0.200796
Carbohydrates   0.781539   0.461213     0.270977  0.200796       1.000000
Dietary Fiber   0.538894   0.580837     0.435575  0.694389       0.224577
Sugars          0.259598  -0.115446    -0.135518 -0.426536       0.762362
Protein         0.787847   0.807773     0.561561  0.869802       0.352122

                Dietary Fiber    Sugars   Protein
Calories             0.538894  0.259598  0.787847
Total Fat            0.580837 -0.115446  0.807773
Cholesterol          0.435575 -0.135518  0.561561
Sodium               0.694389 -0.426536  0.869802
Carbohydrates        0.224577  0.762362  0.352122
Dietary Fiber        1.000000 -0.295178  0.641345
Sugars              -0.295178  1.000000 -0.179940
Protein              0.641345 -0.179940  1.000000
```

```
In [50]: sns.heatmap(cor,square=True)
```

Out[50]: <matplotlib.axes._subplots.AxesSubplot at 0x4af5738ac8>

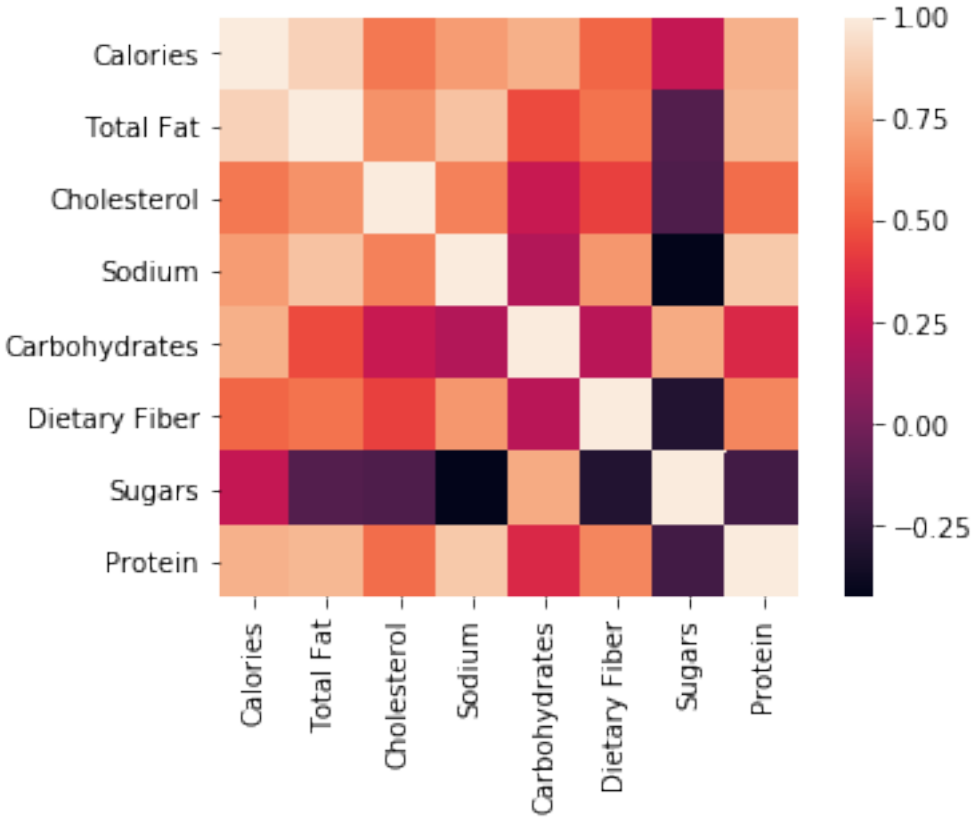

From corrplot we can see total fat,carbohydrates and protein has high positive correlation

4 Split data into train and test

```
In [13]: import sklearn
         from sklearn.model_selection import train_test_split
         X = dataset.loc[:,['Total Fat', 'Carbohydrates','Protein']]
         y = dataset.loc[:,'Calories']
         y.head()
         X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.30, random_state
```

```
In [14]: X_train.shape , X_test.shape, y_train.shape, y_test.shape
```

```
Out[14]: ((182, 3), (78, 3), (182,), (78,))
```

5 Linear Regression Model

```
In [15]: ##Sklearn
         # import model
```

```
from sklearn.linear_model import LinearRegression

# instantiate
linreg = LinearRegression()

# fit the model to the training data (learn the coefficients)
linreg.fit(X_train, y_train)
```

```
Out[15]: LinearRegression(copy_X=True, fit_intercept=True, n_jobs=1, normalize=False)
```

Intercept and Coefficients

```
In [16]: # print the intercept and coefficients
         print(linreg.intercept_)
         print(linreg.coef_)
```

```
-2.227172524597677
[9.00479505 3.99364571 4.05300657]
```

6 Predictions

```
In [17]: # make predictions on the testing set
         y_pred = linreg.predict(X_test)
```

```
In [18]: from sklearn import metrics

         print(np.sqrt(metrics.mean_squared_error(y_test, y_pred)))
```

```
5.784963176081385
```

7 Accuracy

R^2 Value

```
In [20]: linreg.score(X,y)
```

```
Out[20]: 0.9994716750523599
```

~100% accuracy.! R^2 of 0.999(1) means that the independent variables(Total.Fat, Protein, Carbohydrates) are able to explain almost 100% of variance in Calories.